


Standard benchmarks live on the diagonal

Truth and entity frequency agree there, so a probe reading only “common means true” scores the same as one reading truth. The off diagonal is where they disagree, and the only place the question can be asked.



- **Diagonal — where probes are trained**
What the field builds and reports on. Truth and frequency are perfectly correlated here, so both strategies look identical.
- **Off diagonal — where they are tested**
Frequency now points the wrong way. A probe reading truth survives; one reading frequency inverts.

What 18 models show
In distribution, every probe scores between **0.76** and **1.00**. Read honestly on the off diagonal, the same probes span **0.08** to **0.89**.